

## UNIT - 4

DATE: \_\_\_\_\_  
PAGE \_\_\_\_\_

### Transaction Processing System and Concurrency Control

#### Transaction

It is a logical unit of program execution (set of operations) of a database item.

- \* Every transaction in case of failure ~~can be~~ is restarted but not resumed.

Begin transaction

{ set of read / write operations + ALU operations

End transaction

Commit

#### Basic operations of Transaction

1. Read-item ( $x$ )
2. Write-item ( $x$ )

Eg:  $A = 1000 \quad B = 500$

Transfer 50 Rs from A to B.

Read item (A)

$$A = A - 50$$

$$A \rightarrow 1000$$

$$A \rightarrow 950 \quad (\text{ALU})$$

Write (A)  $A \rightarrow 950$

Read-item (B)  $B \rightarrow 500$

$B = B + 50$   $B \rightarrow 550$

Write (B)  $B \rightarrow 550$

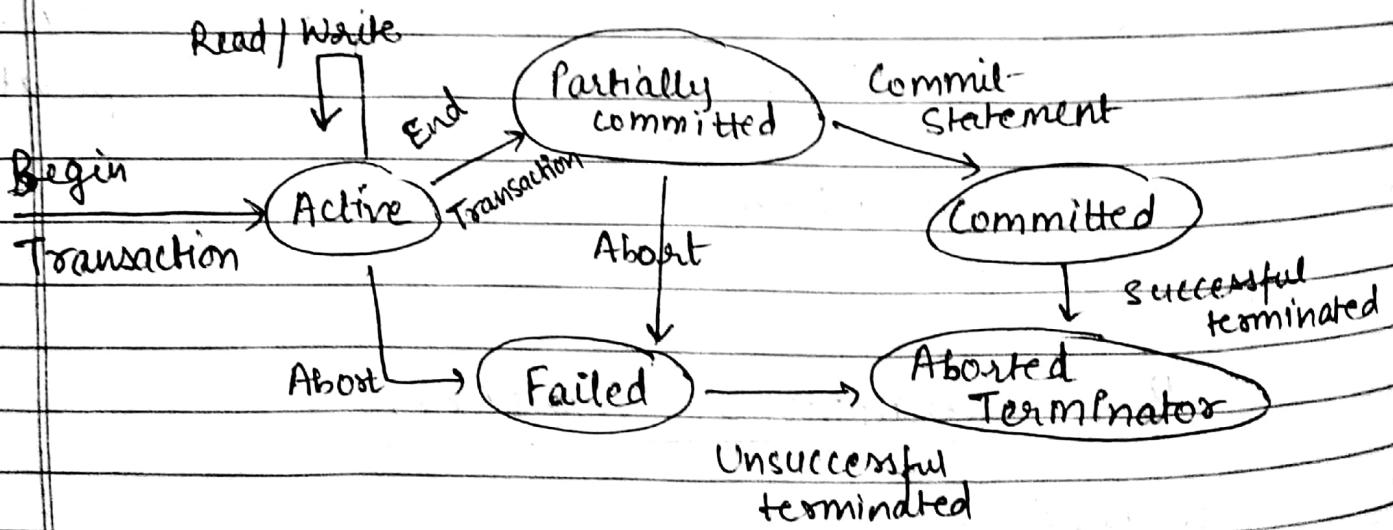
Commit

changes are made in the database

## ACID Properties of a Transaction

- Atomicity
- Consistency
- Isolation
- Durability

## Transaction States



## Schedules (Histories) of Transactions

Suppose  $S_1$  (schedule 1) and  $S_2$  (schedule 2)

①  $S_1 = r_1(x), w_1(x), r_1(y), w_1(y), r_2(x), w_2(x)$

Time

<u>T<sub>1</sub></u>	read(x)	80
	$x = x - N$	75
	write(x)	75
	read(y)	100
	$y = y + N$	105
	write(y)	105

T<sub>2</sub>

read(n)	75
<del><math>n = x + m</math></del>	79
read(x)	75
<del><math>n = x + m</math></del>	79
write(n)	79

(T<sub>1</sub> followed by T<sub>2</sub>)

OR

(2)  $S_2 = r_2(x), w_2(x), r_1(x) w_1(n) r_1(y) w_2(y)$

T<sub>2</sub>.

read(n)	80
$n = x + m$	84
write(x)	84

T<sub>1</sub>.

read(n)	84
$n = x - N$	79
write(n)	79
read(y)	100
$y = y + N$	105
write(y)	105

(T<sub>2</sub> followed by T<sub>1</sub>)

NOTE: The above two schedules: S<sub>1</sub> and S<sub>2</sub> are serial schedules.

(3)

$S_3: r_1(x), r_2(n), w_1(x), r_1(y), w_2(n), w_1(y)$

	$T_1$		$T_2$
read ( $n$ )	80		
$n = n - N$	75		
time			
write ( $n$ )	84		
read ( $y$ )	100		
$y = y + N$	105		
↓			
write ( $y$ )	105		
			write ( $x$ ) 84

(4)

$S_4: r_1(n), w_1(n), r_2(n), w_2(n), r_1(y), w_1(y)$

	$T_1$		$T_2$
read ( $n$ )	80		
$n = n - N$	75		
time			
write ( $n$ )	75		
↓			
read ( $y$ )	100		
$y = y + N$	105		
write ( $y$ )	105		
			read ( $x$ ) 75
			$n = n + m$ 79
			write ( $x$ ) 79

NOTE :  $S_3$  and  $S_4$  are non-~~schedule~~<sup>serial</sup> schedules.  
 Execution is done in an interleaved way  
 (In case of concurrent execution of transactions)

GOOD WRITE

\* When same database elements is accessed by two other / different transactions, then anomalies occur.

Eg: S3 in which  $x$  is accessed by  $T_1$  and  $T_2$ .

Q. Why Concurrency control is needed?

There are three types of problems:

① Lost Update Problem

In S3, the final value of  $x$  is incorrect because  $T_2$  reads the value of  $x$  before  $T_1$  changes it in the database. Hence, updated value from  $T_1$  is lost.

~~Q~~ ② Dirty Read Problem (Temporary Update)

If in S4, transaction  $T_1$  fails and must change the value of  $x$  back to its old value; meanwhile  $T_2$  has read the temporary incorrect value of  $x$ .

S4	$T_1$	$T_2$
	read( $x$ )	read( $x$ )
	80	75
	$x = x - N$	$x = x + m$
	75	79
	write( $x$ )	write( $x$ )
time		
failure	→	
	read( $y$ )	read( $y$ )
	$y = y + N$	$y = y + m$
	GOOD WRITE	GOOD WRITE
	write( $y$ )	write( $y$ )

③ Incorrect Summary Problem.

T<sub>1</sub>

```
read(x);
n = x - N;
write(n);
```

```
read(y);
y = y + N;
write(y);
```

T<sub>2</sub>

```
sum = 0;
read(A);
sum = sum + A;
```

```
read(x);
sum = sum + x;
read(y);
sum = sum + y;
```

T<sub>2</sub> reads the value of  $x$  after  $N$  ~~seats have~~  
 has been subtracted from it but it reads  
 the value of  $y$  ~~after~~ before  $N$  has been  
 added to it. Hence, a wrong  
 summary is the result (off by  $N$ )

# Concurrency Control Technique (CCT)

① Locking  
Based Control

② Timestamp Ordering  
Based Control

③ Multiversion control

④ Validation  
Based Protocol  
(Optimistic)

Binary  
Lock

Shared /  
Exclusive Lock

Two Phase  
Locking  
(2PL)

## Binary Lock

$\text{lock}(A)$  → Here, A is any data item used in a transaction T

$\text{unlock}(A)$  → release A so that the other transactions can take A.

## Shared / Exclusive Lock

Shared - read only (also called share-locked)  
Exclusive - read and write both (also called exclusive-locked)

3 operations :-

1.  $\text{read-lock}(A)$  → shared-lock
2.  $\text{write-lock}(A)$  → exclusive-lock
3.  $\text{unlock}(A)$  → Both shared lock and exclusive lock

89

$T_1$   
 read(A)  
 read(B)  
 $B = B + A$   
 write(B)



read-lock(A)  
 read(A)  
 unlock(A)  
 writer-lock(B)  
 read(B)  
 $B = B + A$   
 write(B)  
 unlock(B)

### Compatible Graph.

$T_1$	$T_2$	S	X
	S	True	False
	X	False	False

Q1  $T_1$ : Account X transfers Rs 50 to account Y

$T_1$   
 200  
 read(x)  
 $x = x - 50$   
 150  
 write(x) →  
 100  
 read(y)  
 $y = y + 50$   
 150  
 write(y)

$T_1'$  (Lock S/X)  
 write-lock(x)  
 read(y)  
 $y = x - 50$   
 write(x)  
 unlock(x)  
 write-lock(y)  
 read(y)  
 $y = y + 50$   
 write(y)  
 unlock(y)

$$\begin{aligned} x &= 150 \\ y &= 150 \\ \text{Sum} &= 300 \end{aligned}$$

Serial Order

$T_1 \rightarrow T_2'$

OR

$T_2 \rightarrow T_1'$

DATE: \_\_\_\_\_  
PAGE: \_\_\_\_\_

Q: Go simply add x and y and display their sum in sum variable

$T_2$

150 read (x)  
150 read (y)  
300 sum = x + y  
300 display (sum)

$T_2'$

read-lock (x) 200  
read (x) 200  
unlock (x)  
read-lock (y) 100  
read (y) 100  
unlock (y)  
sum = x + y 300  
display (sum) 300

$$\begin{aligned} x &= 200 \\ y &= 100 \\ \text{sum} &= 300 \end{aligned}$$

\* interleaved transactions  
next page →

## Two Phase Locking

The two phases are :-

1. Growing / Expanding phase :- new locks on the items can be acquired but none can be released
2. Shrinking phase :- the existing locks can be released but no new locks can be acquired

Eg:

read (A)  
read (B)  
write (B)  
read (C)  
write (C)



read-lock (A)  
read (A)  
write-lock (B)  
read (B)  
write (B)  
write-lock (C)  
read (C)  
write (C)  
unlock (C)  
unlock (A)  
unlock (B)

~~$\theta^1$  and  $\theta^2$~~

## Interleaved way

$T_1'$

$T_2'$

Write-lock (n)	200
read (n)	200
$n = n - 50$	150
Write (n)	150
unlock (n)	

read-lock (y)	100
read (y)	100

read-lock (n)	150
read (n)	150
unlock (n)	

sum = n + y	250
display (sum)	250

write-lock (y)	100
read (y)	100
$y = y + 50$	150
write (y)	150
unlock (y)	

$$n = 150$$

$$y = 150$$

$$\text{sum} = 250$$

Here, incorrect sum is calculated because  $T_1$  unlocked data from  $x$  too early, as a result  $T_2$  shows an inconsistent state.

## Rules for time stamp protocols

1) If a transaction  $T_i$  issues read operation

- a) If  $TS(T_i) \geq WTS(Q)$ , the read operation is executed and  $RTS(Q)$  is set to  $\max(RTS(Q), TS(T_i))$ .  
 → (a, b, c variables or data items)

b) Otherwise c

If  $TS(T_i) < WTS(Q)$ , the read operation is rejected and transaction  $T_i$  is restarted with a new time stamp value.

2) If a transaction  $T_i$  issues write operation

- a) If  $TS(T_i) \geq WTS(Q)$ , the write operation is executed and  $WTS(Q)$  is set to  $TS(T_i)$ .

- b) If  $TS(T_i) < RTS(Q)$ , the write operation is rejected and transaction  $T_i$  is restarted with new time stamp value.

Example:

Step	$T_1$ (200)	$T_2$ (150)	$T_3$ (175)	Q	b	c
				$RT(a) = 0$ $WT(a) = 0$	$RT(b) = 0$ $WT(b) = 0$	$RT(c) = 0$ $WT(c) = 0$
1	Read(b)					$RT(b) = 200$
2		Read(a)				$RT(a) = 150$
3			Read(c)			$RT(c) = 175$
4	Write(b)					$WT(b) = 200$
5	Write(a)					$WT(a) = 200$
6		Write(c)				$WT(c) = 150$
7			Write(a)			

TS(T<sub>1</sub>) < WTS(b) (False)

Step 1 : TS(T<sub>1</sub>) = 200 > 0 (Yes)

$$RTS(b) = \max(0, 200) = 200$$

Step 2 : TS(T<sub>2</sub>) = 150 > 0 (Yes)

Read operation permit

$$RTS(a) = \max(0, 150) = 150$$

Step 3 : TS(T<sub>3</sub>) = 175 > 0 (Yes)

Read operation permit

$$RTS(c) = 175$$

Step 4 :

TS(T<sub>1</sub>) ~~> RTS~~ >= WTS(b)

$$200 >= 0 \quad (\text{True})$$

Write operation permits

$$WTS(b) = TS(T_1) = 200$$

Step 5 : TS(T<sub>1</sub>) >= WTS(a)

$$200 >= 0 \quad (\text{True})$$

Write operation permit

$$WTS(a) = 200$$

Step 6 : TS(T<sub>2</sub>) >= WTS(c)

$$150 >= 0 \quad (\text{True})$$

Write operation permits

$$WTS(c) = 150$$

Step 7 : TS(T<sub>3</sub>) < WTS(a) {175 < 200}

Write operation is rejected

## Validation or optimistic concurrency control Techniques.

- It assumes that conflicts between the transactions are rare.
- This method does not require locking and hence no deadlock occurs.
- This method checks for ~~conflicts~~ just before commit.

### Basic terminologies used

Read set ( $T_j$ ) : set of items read by  $T_j$  (transaction  $j$ )

Write set ( $T_j$ ) : set of items written by  $T_j$

$Start(T_j)$  ] 2 local variables of  $T_j$  containing  
 $Finj(T_j)$  timestamp value that are needed  
for validation phase

$Start(T_j)$  → when transaction  $T_j$  started

$Finish(T_j)$  → when transaction  $T_j$  finished its read phase

$TS(T_j)$  : Unique timestamp value assigned to  $T_j$  only after write phase if the validation succeeds.

### Three phases of validation / optimistic CC Technique.

1. Read phase (R)
2. Validation phase (V)
3. Write phase (W)

R	V	W
---	---	---

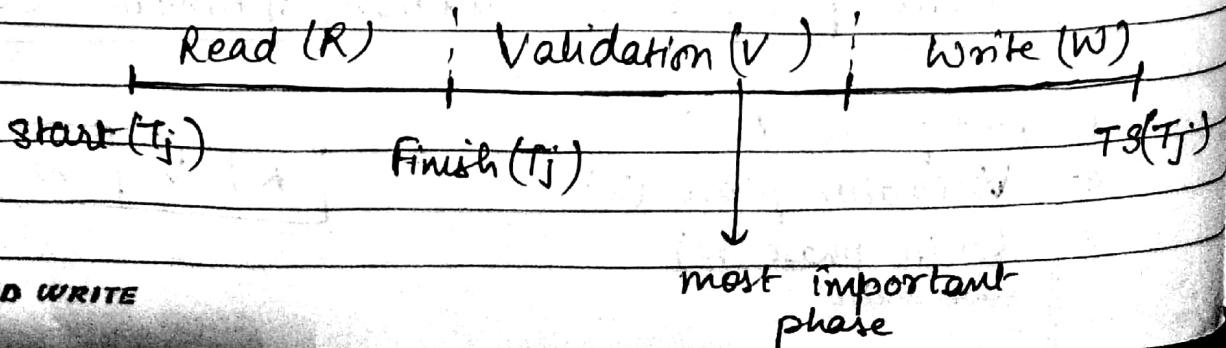
1. Read Phase → In this phase, a transaction reads the data item from the database, performs computation and writes all results into local or private variables i.e. writes locally. These variables are not written into the database.

2. Validation Phase (to check consistency) → This phase ensures that for those data items that were read and modified (updated) in the read phase will not cause the database to become inconsistent.

If any inconsistency occurs, then transaction will be rolled ~~back~~ back and restarted with new timestamp value.

3. Write phase → If the validation phase succeeds, the transaction updates (writes) are applied permanently to the database. Otherwise, transactions are restarted with new timestamp value.

Structure of transaction  $T_j$

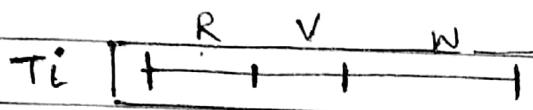


⇒ Validation phase can be completed successfully if atleast one of the following 3 conditions is satisfied:

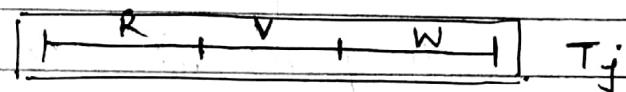
Condition 1:

$$\boxed{TS(T_i) < Start(T_j)}$$

All the older transactions ( $T_i$ ) must have completed their write phase before  $T_j$  started its read phase.



$TS(T_i)$



$Start(T_j)$

They must execute in serial order ∴ no interference.

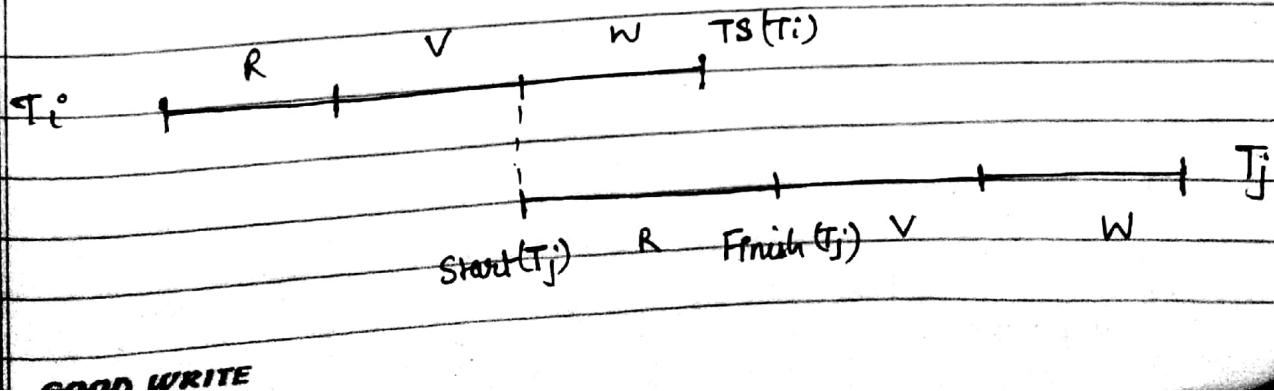
Condition 2:

$$\boxed{Start(T_j) < TS(T_i) < Finish(T_j)}$$

and

$$\boxed{\text{Writeset}(T_i) \cap \text{Read set}(T_j) = \emptyset}$$

i.e. no common data elements



good WRITE

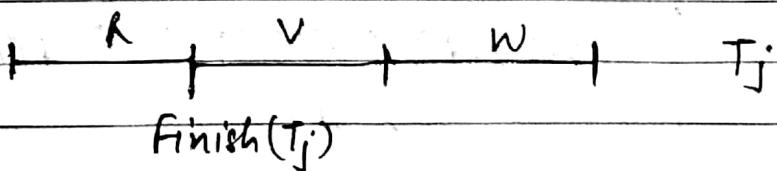
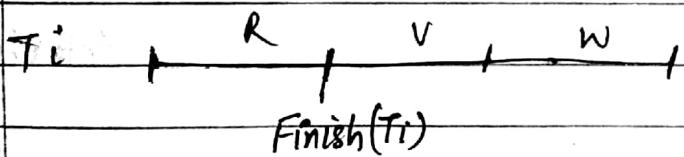
$T_i$  completes its write phase before  $T_j$  starts its validation phase and the write set of  $T_i$  does not intersect with the readset of  $T_j$ .

### Condition 3

$\text{Finish}(T_i) < \text{Finish}(T_j)$   
and

$$\text{Write set } T_i \cap (\text{readset } (T_j) \cup \text{Write set } (T_j)) = \emptyset$$

- $T_i$  completes its ~~read~~ read phase before  $T_j$  completes its read phase.
- The write set of  $T_i$  does not intersect with readset or write set of  $T_j$ .



## Backup and Recovery

### System Log

Start (T)

Write (T, X, old-value, new-value)

Read (T, X)

Commit (T)

Rollback (T)

### Recovery Techniques

#### 1. Deferred Update Recovery (NO-UNDO / REDO technique)

- During transaction execution, updates are written to the log file.
- After transaction reaches the commit point, the log file is force-written to disk, then updates are recorded in the database.
- A transaction cannot change the database on disk until it reaches its commit point.

Example : T<sub>1</sub>: Ra Rd Wd C  
T<sub>2</sub>: Rb Wb Rd Wd C

### Log File

Start (T<sub>1</sub>)

Read (T<sub>1</sub>, a)

Read (T<sub>1</sub>, d)

Write (T<sub>1</sub>, d, old, new)

Commit ( $T_1$ )

Start ( $T_2$ )

Read ( $T_2, b$ )

Write ( $T_2, b, \text{old}, \text{new}$ )

System crash.

→ Redo

~~# fail~~ Since  $T_1$  has reached commit,  
 $\therefore$  changes in  $T_1$  are written to disk.

However,  $T_2$  did not commit, hence its changes were not written to disk. To recover;

\* UNDO (Backward Recovery)  $\rightarrow$  It is applied to a single operation rather than to whole transaction.

\* REDO (Forward Recovery)  $\rightarrow$  operations in a certain transaction must be redone to ensure that all the operations of a committed transaction have been applied successfully to the database.